

Fictora Content Operator Guide

Fictora episode production — operator guide

Audience: Mihir, Tejas, and anyone filming ep1 on prod with Cursor + this repo.
Updated: September 2026 · **Tooling:** `uv run fictora-produce` (not raw API calls).

This is the day-to-day playbook. Craft rules and API detail live in [runbook.md](#) and [api-map.md](#). First-time machine setup: [Content-Producer-Setup-Handout.md](#).

What you are doing

You produce **9:16 vertical ~15s takes** on Fictora’s hosted Drama API. You work in two places:

Place	What happens
Cursor (this repo open)	The agent runs one paid API step per turn; you say aligned , yes , or deviation .
Desk folder (~/Downloads/documents/YYYY-MM-DD-series-slug/)	You open plates , boards , and takes and approve before money moves.

You never need the private `fictora-drama` repo. Never paste the API token into chat.

One-time setup (each person, ~10 min)

1. Clone `DreamFlux-Workspace/fictora-content-creation` and run `uv sync`.
2. Copy `.env.example` → `.env` and set **FICTORA_DRAMA_GENERATION_SERVICE_TOKEN** (from engineering).
3. Open the **repo root** in Cursor (so the **episode-production** skill loads).
4. Optional check (no spend): `uv run python scripts/smoke_live.py --phase read`
5. For local yellow captions after filming, install **ffmpeg with libass** on macOS: `brew install ffmpeg-full` (agent uses `/opt/homebrew/opt/ffmpeg-full/bin/ffmpeg`).

Start a new series (paste into Cursor)

Use a **new series name** and a **dated slug** (folder name is automatic, e.g. 2026-09-22-tram-not-tonight).

Read ``.cursor/skills/episode-production/SKILL.md`` and follow it exactly.

Goal: Produce episode 1 on prod Drama API – new series.

Story (ep1):

- Vertical 15s, modern-dark-fantasy preset, minimax-h3.
- [Setting, beat, one speaking line, cast notes, no gore.]

Desk:

- ``uv run fictora-produce start`` with series "... " and your premise.
- Use ``--draft-episodes 4``, ``--clip-seconds 15``, ``--cut-tempo one_shot``, ``--caption-style house``.
- Do NOT pass ``--api-captions`` (default: raw clip on server, captions on laptop).

Workflow:

- One ``fictora-produce step`` per turn until I approve gates.
- After boards, estimate → I say yes → ``step --confirm-spend``.
- When raw MP4 lands, burn house captions with ``scripts/burn_house_captions.sh`` and open the captioned take.

Start with ``start``, print desk path and phase, wait for my "aligned" before the first paid step.

Prod-validated example (copy premise flags from here): [examples/tram-not-tonight-ep1.md](#).

Your gates (human yes before money)

Order	You open	You say	Agent runs after your yes
1	—	aligned	First step (draft plan)
2	ep01/plates/	plates ok / next	approve -- gate plates then step
3	Script in chat or ep01/api/03_spine.json beats	script ok / next	approve -- gate script then step (boards)
4	ep01/boards/ + exposure note in chat	board ok / next	approve -- gate board –

Order	You open	You say	Agent runs after your yes
			-path ... (add - -accept-dim if board is dark night interior) step -- confirm- spend
5	Estimate in chat (~\$1.20 take)	yes	
6	ep01/takes/*captioned*.mp4	use or deviation + cause	Re-board or re- film only with a named cause

Silence is not approval. Each gate needs a fresh yes.

Phase cheat sheet (what “next” means)

Phase	Meaning
new	Desk ready; waiting for aligned → draft
wait_plates	Review cast stills
wait_script	Review ep1 lines on spine
wait_board	Review storyboard; check dim warning (~25% luma floor)
wait_spend	Confirm yes before ~\$1.20 take
complete	Raw MP4 on disk → local captions + QC

Check anytime:

```
uv run fictora-produce status --desk ~/Downloads/documents/<your-desk-
folder>
```

Commands (you or the agent)

Replace <desk> with the full path printed at start.

Create desk (no API spend yet)

```
uv run fictora-produce start \
  --series "Series Title" \
  --prompt "Premise..." \
  --band 15s \
  --preset-id modern-dark-fantasy \
  --video-lane minimax-h3 \
  --draft-episodes 4 \
  --clip-seconds 15 \
  --cut-tempo one_shot \
```

```

--caption-style house

# One automated API block per invocation
uv run fictora-produce step --desk <desk>

# Human gates
uv run fictora-produce approve --desk <desk> --gate plates
uv run fictora-produce approve --desk <desk> --gate script
uv run fictora-produce approve --desk <desk> --gate board \
    --path <desk>/ep01/boards/board-ep01-t1-1-v1.png
# If exposure was below dim floor and you accept the dark board:
#     ... --accept-dim

# Film take (after you said yes to estimate)
uv run fictora-produce step --desk <desk> --confirm-spend

```

Do not use --draft-episodes 5 — prod plans four episodes; five causes plan errors.

Do not use --api-captions unless engineering asks — it slows post and waits on server burn-in.

After the raw take (local house captions)

Default config: **api_captions: false**. Filming finishes when the **raw** scene MP4 is downloaded (often within a minute after MiniMax returns). You do **not** wait for server caption burn.

File	Role
ep01/takes/take-ep01-t1-raw-v1.mp4	Silent of burned captions; use for timing
ep01/api/17_raw_scene_clips.json	Same clip URL
ep01/takes/take-ep01-t1-captioned-v1.mp4	Ship candidate after local burn

Agent workflow ([local-captions.md](#)):

1. Line text from spine / script gate (e.g. “*Not tonight.*”).
2. ffmpeg **silencedetect** on the raw take → anchor captions to **speech span**, not long pauses.
3. Write take-ep01-t1-house.ass (yellow #F5D547, black outline, no box).
4. Burn:

```

cd /path/to/fictora-content-creation
FFMPEG=/opt/homebrew/opt/ffmpeg-full/bin/ffmpeg \
    bash scripts/burn_house_captions.sh \
    <desk>/ep01/takes/take-ep01-t1-raw-v1.mp4 \
    <desk>/ep01/takes/take-ep01-t1-house.ass \
    <desk>/ep01/takes/take-ep01-t1-captioned-v1.mp4
open <desk>/ep01/takes/take-ep01-t1-captioned-v1.mp4

```

Scrub to the spoken line; captions should flicker word-by-word in the lower centre band.

Desk layout

2026-09-22-my-series/	
production.json	← phase, spine_id, session (agent)
production.config.json	← 15s, one_shot, house, api_captions false
QUEUE.md	← what episode is waiting on
ep01/	
plates/	← gate 1
boards/	← gate 2
takes/	← raw + captioned video
api/	← JSON log (support / debugging)
run-notes.md	← ledger (agent updates)
brief.md	

Spend (H3 lane, approximate)

Unit	~USD
Cast plates (batch)	\$0.30
Storyboard	\$0.30
15s take	\$1.20
First ep envelope	~\$4.50

Stop and escalate if spend goes past **2×** the envelope without a written reason.

A **second** take needs a **named cause** (brief/board issue), not “try again.”

Troubleshooting

Symptom	What to do
Empty plates/ or boards/ right after step	Pull latest main; re-run step on same desk if phase allows, or ask agent to download from ep01/api/06_*cast_terminal*.json / 09_*boards_terminal*.json.
Network error while polling	Same desk, run step again — job may still be running server-side.
Stuck at 50% on video with local captions	Often post-prod you can ignore; with api_captions: false the harness should already stop at raw clip. Check for 17_raw_scene_clips.json.

Symptom	What to do
Board “dim” warning (~14–24% luma)	Night interiors are common; approve with --accept-dim if the mosaic looks right on screen.
Plan fails on episode 5	Set draft_episode_count to 4 only.
Token / 503 / billing	Engineering — do not rotate token in chat.

Video recovery (agent, same desk): `fictora-produce cancel-job → retry-video → step --confirm-spend` (see episode-production skill).

Rules of the road

- Open every gate file before you say ok.
- One **paid** enrol per agent turn unless you explicitly ask to catch up.
- Never commit `.env` or share the service token.
- Never approve plates, script, and board in one message without looking.

References

Doc	Use
examples/tram-not-tonight-ep1.md	Copy-paste start command + prod outcomes
Content-Producer-Setup-Handout.md	Setup + PDF source
local-captions.md	Caption timing detail
runbook.md	Craft, sound, caption recipe
<code>.cursor/skills/episode-production/SKILL.md</code>	Agent contract

Support: Token, outages, billing → DreamFlux engineering. Day-to-day gates → this guide + Cursor agent.